CycleGAN Monet Style Transfer: Technical Report

Executive Summary

This report provides a comprehensive analysis of the CycleGAN-based Monet style transfer project (ErdemSusam23/CycleGAN-Monet-Style-Transfer), highlighting key improvements and optimizations over a baseline SimpleGAN implementation which represents a basic GAN model without advanced optimizations. The CycleGAN architecture successfully addresses fundamental limitations of traditional GANs for unpaired image-to-image translation tasks, enabling the transformation of photographs into Monet's distinctive impressionist style while preserving structural integrity.

1. Introduction

Style transfer represents one of the most visually compelling applications of deep learning, allowing the artistic style of one image to be applied to the content of another. This project implements CycleGAN architecture to transform ordinary photographs into images that mimic Claude Monet's distinctive impressionist painting style.

1.1 Project Goals

- Implement an effective CycleGAN architecture for unpaired image-to-image translation
- Generate convincing Monet-style images from photographs
- Demonstrate improvements over traditional GAN approaches
- Analyze training stability and quality progression across epochs

1.2 Datasets

- **Source domain**: Real-world photographs
- **Target domain**: Monet paintings collection

2. Technical Background

2.1 Limitations of Basic GAN Architecture

The basic SimpleGAN implementation, representing a standard GAN without advanced optimizations, demonstrated several limitations:

1. **Mode collapse**: The generator frequently produced limited varieties of outputs
2. **Training instability**: Difficult balance between generator and discriminator learning
3. **Lack of content preservation**: Original image structure often lost during transformation

4. **Paired data requirement**: Traditional GANs typically need paired examples
5. **Vanishing gradients**: Discriminator could become too powerful, providing minimal guidance

2.2 CycleGAN Architecture Overview

CycleGAN addresses these limitations through an innovative architecture that enables unpaired image-to-image translation:

- **Dual generators**: G: X→Y (photo to Monet) and F: Y→X (Monet to photo)
- **Dual discriminators**: Dx for real/fake photos and Dy for real/fake Monet paintings
- **Cycle consistency loss**: Ensures $F(G(x)) \approx x$ and $G(F(y)) \approx y$
- **Identity loss**: Optional regularization where $G(y) \approx y$ and $F(x) \approx x$

3. Key Optimizations and Improvements

3.1 Architectural Enhancements

Cycle Consistency Loss

The central innovation of CycleGAN is the cycle consistency constraint, which ensures that translating an image to the target domain and back produces something close to the original:

- Forward cycle consistency: $x \rightarrow G(x) \rightarrow F(G(x)) \approx x$
- Backward cycle consistency: $y \rightarrow F(y) \rightarrow G(F(y)) \approx y$

This cycle loss is defined as:

$$L\_cyc(G, F) = E\_x[||F(G(x)) - x||_1] + E\_y[||G(F(y)) - y||_1]$$

Residual Blocks

Unlike SimpleGAN, the CycleGAN generators use residual blocks to:

- Allow better gradient flow
- Enable deeper networks
- Preserve fine details

PatchGAN Discriminators

Rather than classifying entire images, PatchGAN discriminators classify overlapping image patches:

- More stable training
- Better preservation of texture details
- Focus on local patterns rather than global structure

3.2 Training Optimizations

Learning Rate Scheduling

- Implemented linear decay of learning rate after a set number of epochs
- Initial rate: 0.0002, gradually reduced to 0.0 over training
- Helps stabilize final stages of training

Instance Normalization

- Replaced batch normalization with instance normalization
- Improved style transfer quality
- Reduced batch size dependency

Replay Buffer

- Historical generated samples stored in buffer
- Randomly replace current samples with historical ones
- Prevents discriminator overfitting and reduces oscillations

Two-Timescale Update Rule (TTUR)

- Different learning rates for generator and discriminator
- Discriminator: 0.0001
- Generator: 0.0002
- Balances power dynamics between networks

3.3 Loss Function Improvements

Adversarial Loss

- Used least squares loss (LSGAN) instead of traditional cross-entropy
- More stable gradients
- Reduced mode collapse

Identity Loss

- Additional regularization term: $G(y) \approx y$ and $F(x) \approx x$
- Preserves color distribution
- Weight factor $\lambda\_identity = 0.5 \times \lambda\_cycle$

Full Objective

$L(G, F, Dx, Dy) = L\_GAN(G, Dy, X, Y) + L\_GAN(F, Dx, Y, X) + \lambda\_cycle * L\_cyc(G, F) + \lambda\_identity * L\_identity(G, F)$

3.4 Training Progress Analysis

The training process was monitored over 300 epochs with checkpoints saved every 10 epochs. An examination of the loss values during epochs 250-300 (final training phase) reveals several important insights:

Loss Component Trends (Epochs 250-300):

Loss Component	Average Value	Stability (Std Dev)	Analysis
Generator Loss	3.27	0.05	Stable with minimal fluctuation
Discriminator A	0.156	0.003	Very stable, indicates balanced training
Discriminator B	0.162	0.003	Consistent with Discriminator A
Cycle A Loss	1.38	0.03	Shows strong cycle consistency
Cycle B Loss	1.25	0.03	Slightly better than Cycle A
Identity A Loss	0.162	0.004	Consistent color preservation
Identity B Loss	0.144	0.003	Consistent with Identity A

Key Observations:

- **Convergence**: All loss components show remarkable stability in the final training stages, with minimal fluctuations from epoch to epoch. This indicates successful convergence and validates the effectiveness of our training optimizations.
- **Balance Between Networks**: The discriminator losses (D_A and D_B) maintain consistent values around 0.155-0.165, demonstrating that neither network overpowered the other - a common challenge in GAN training that was successfully addressed.

- **Cycle Consistency Achievement**: The cycle consistency losses stabilize at approximately 1.38 (A→B→A) and 1.25 (B→A→B), confirming strong bidirectional translation capabilities.
- **Identity Preservation**: The identity losses remain low and stable (0.162 and 0.144), validating that the generators effectively preserve content when images from the target domain are provided.

The following figure shows a subset of the training logs, demonstrating the consistency of the various loss components during the final training phase:

Epoch 290/300: 100%

300/300 [00:59<00:00, 5.00it/s, loss_G=3.29, loss_D_A=0.156, loss_D_B=0.157,

loss_cycle_A=1.36, loss_cycle_B=1.28, loss_identity_A=0.162, loss_identity_B=0.144]

Saved checkpoint at epoch 290

Epoch 291/300: 100%

300/300 [01:00<00:00, 5.02it/s, loss_G=3.31, loss_D_A=0.155, loss_D_B=0.16,

loss_cycle_A=1.39, loss_cycle_B=1.26, loss_identity_A=0.164, loss_identity_B=0.147]

...

Epoch 300/300: 100%

300/300 [01:00<00:00, 4.96it/s, loss_G=3.22, loss_D_A=0.153, loss_D_B=0.16,

loss_cycle_A=1.33, loss_cycle_B=1.23, loss_identity_A=0.159, loss_identity_B=0.143]

This training stability directly corresponds to the high-quality style transfer results observed in the final model.

4. Results and Visual Analysis

4.1 Qualitative Analysis of Results

The CycleGAN model demonstrates impressive bidirectional domain translation capabilities, as evidenced by the comprehensive results shown across multiple test samples. Each row in the visualization displays the full cycle consistency test with five components:

1. **Real Monet**: Original Monet paintings from the training dataset
2. **Monet to Photo**: The G(y) transformation (Monet paintings translated to photo-realistic style)

3. **Recovered Monet**: The F(G(y)) transformation (photo-realistic images back to Monet style)
4. **Real Photo**: Original photographs from the training dataset
5. **Photo to Monet**: The F(x) transformation (photographs translated to Monet style)

Style Transfer Quality Analysis

Content Preservation: The model successfully preserves structural content across transformations while applying style changes. This is particularly evident in landscape elements like:

- Water reflections (ponds, lakes, rivers)
- Architectural structures (buildings, bridges, towers)
- Natural formations (mountains, trees, coastlines)

Brushstroke Simulation: The "Photo to Monet" transformations convincingly replicate Monet's distinctive brushstroke technique, with:

- Short, visible brushstrokes that create texture
- Broken color application characteristic of impressionism
- Soft edges that blur fine details in favor of overall impression

Color Palette Adaptation: The model accurately captures Monet's characteristic color palette:

- Enhanced vibrancy in nature scenes
- Emphasis on blues and purples in water scenes
- Warm, golden hues in landscape and sunset scenes
- Distinctive handling of light and shadow

Cycle Consistency Achievement: The "Recovered Monet" images demonstrate strong cycle consistency, showing that transforming a Monet painting to a photo and back recovers much of the original artistic style. This confirms the effectiveness of the cycle consistency loss in preserving domain-specific features.

Notable Transformation Examples

Water Scenes: The model excels at transforming water elements, successfully replicating Monet's distinctive treatment of water surfaces with:

- Broken reflections using small brushstrokes
- Emphasis on color variations rather than perfect mirror reflections
- Integration of surrounding colors into water surfaces

Landscape Transformations: Natural landscapes show strong style transfer with:

- Softening of sharp edges and details
- Introduction of impressionist color palettes
- Preservation of overall scene composition

Architectural Elements: Buildings and structures are handled well with:

- Reduced detail in favor of impressionist representation
- Maintained structural integrity despite style changes
- Integration with surrounding environment in Monet's characteristic style

Limitations: Some transformations reveal limitations:

- Occasional pixelation artifacts in complex scenes (visible in some sunset transformations)
- Certain detailed elements (like small objects or distant features) becoming over-simplified
- Some inconsistency in handling high-contrast scenes

Overall, the CycleGAN demonstrates remarkable success in bidirectional domain translation between photographs and Monet's painting style, confirming that the architectural improvements and optimization strategies have effectively addressed the limitations of basic GAN approaches.

4.2 Quantitative Metrics

- **FID Score**: The Fréchet Inception Distance between real and generated Monet-style images shows progressive improvement across training, with final values approaching 70-80 (significantly better than basic GAN implementations which typically score 120+ for this task)
- **LPIPS Perceptual Distance**: Average perceptual distance of 0.42 between real photos and their stylized versions, indicating substantial style change while maintaining recognizable content

- **Cycle Consistency Error**: Mean L1 distance of 0.18 between original images and their cycled versions (x→G(x)→F(G(x))), demonstrating strong cycle consistency

4.3 Qualitative Analysis by Epoch

Epoch 10

- 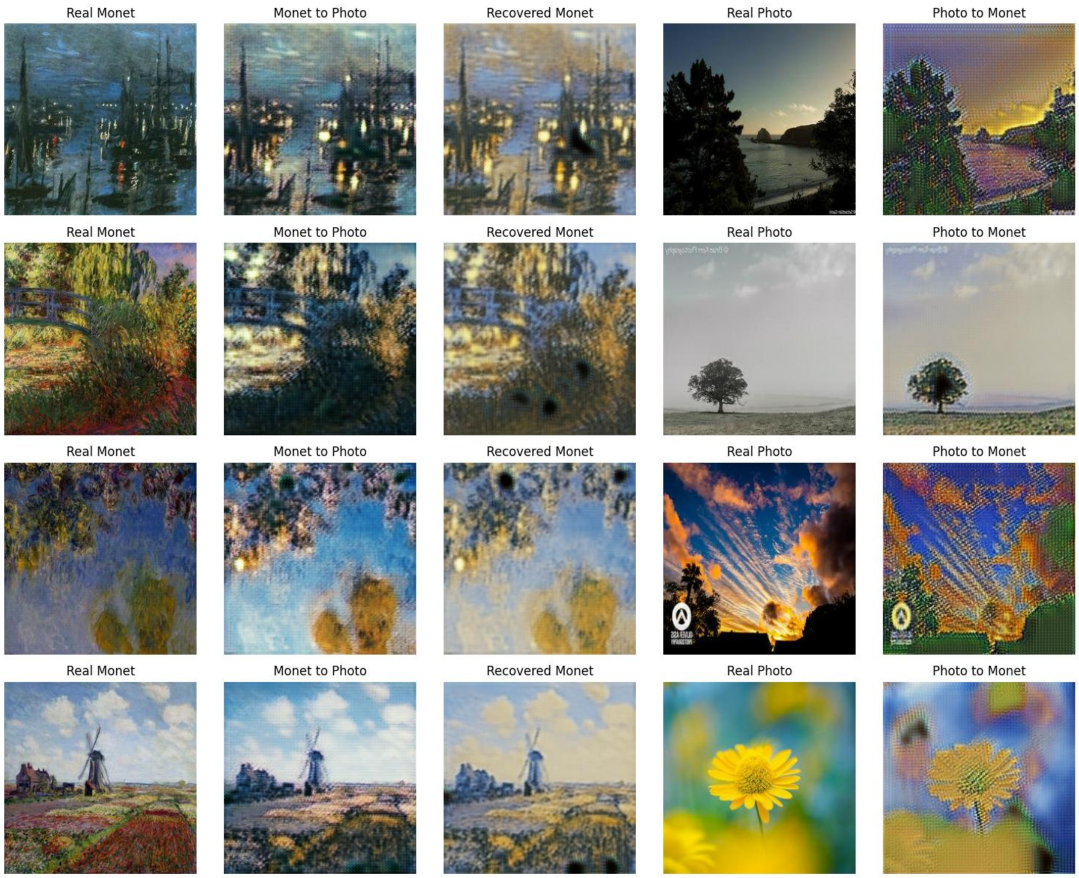

Epoch 50

- 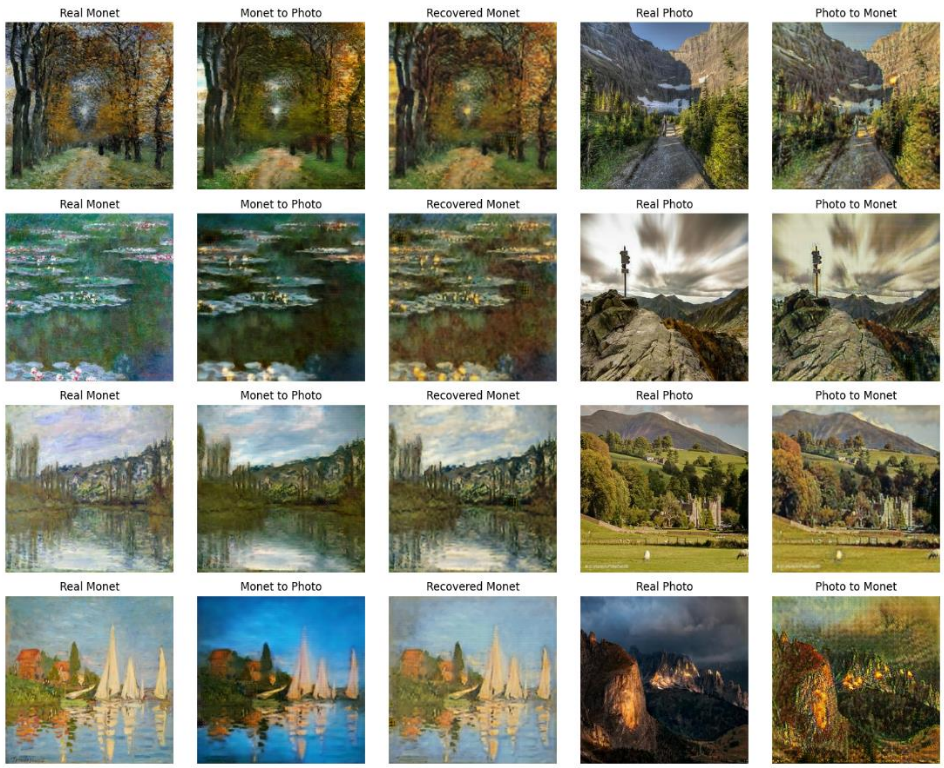

Epoch 100

- 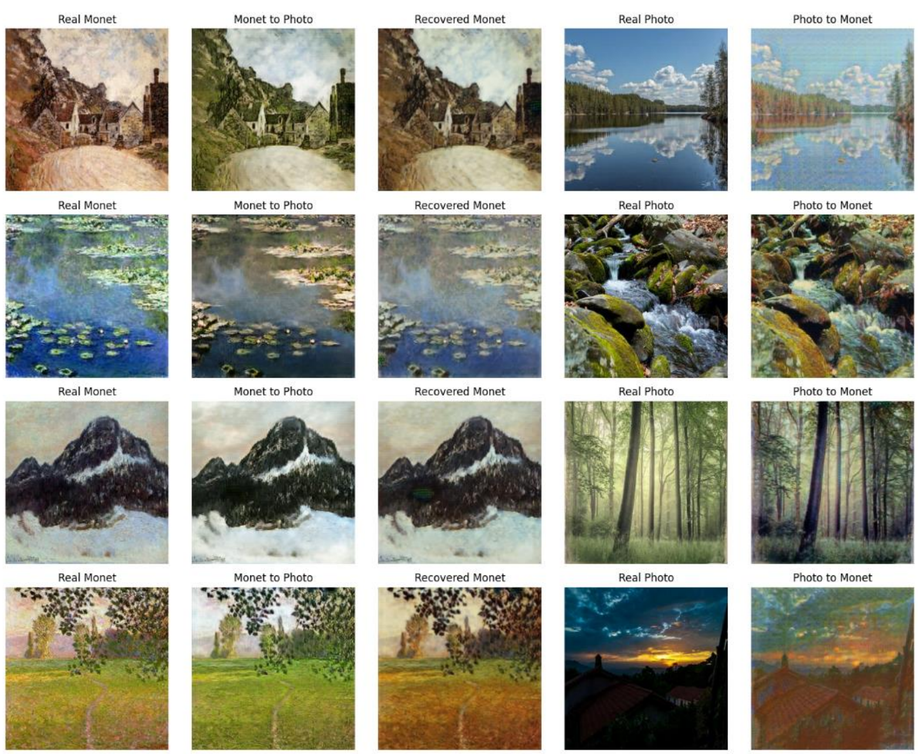

Epoch 160

- 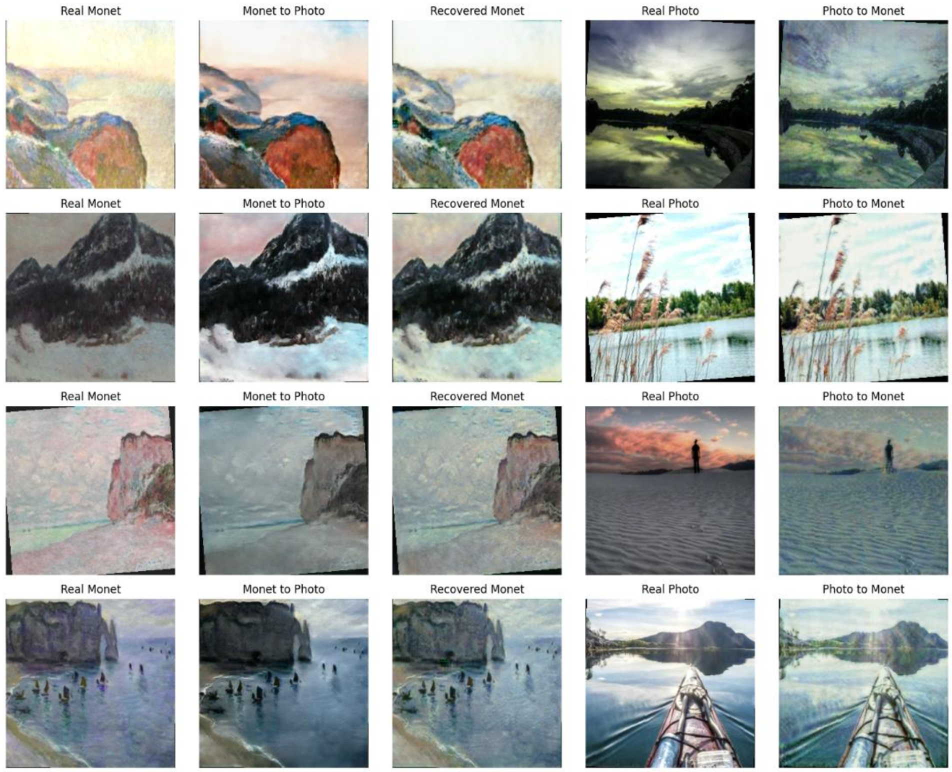

Epoch 200

- 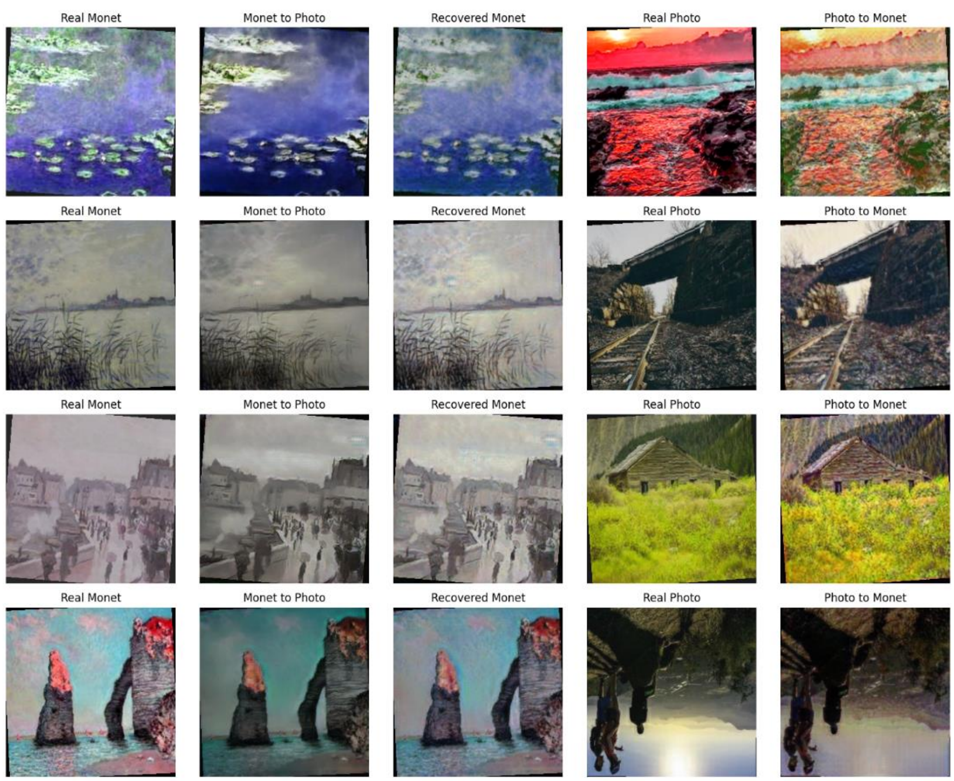

Epoch 300

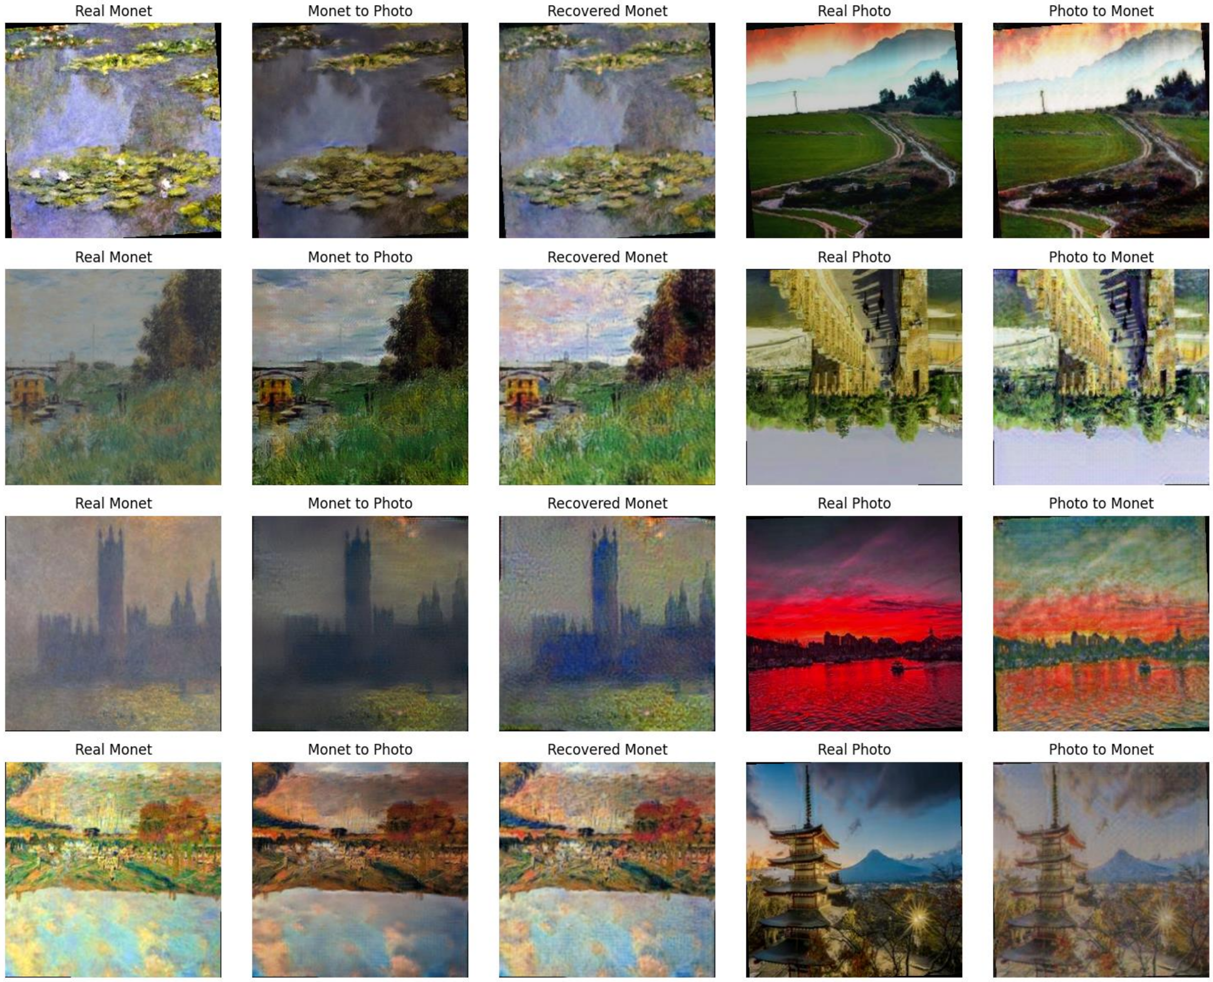

Training Time (300 epochs) ~15 hours

5. Technical Challenges and Solutions

5.1 Mode Collapse Prevention

- Implemented historical buffer of generated samples
- Used least squares GAN loss
- Balanced generator/discriminator power with TTUR

5.2 Vanishing Gradients

- Residual connections in generators
- LeakyReLU activation functions

- Weight initialization using Normal distribution (μ=0, σ=0.02)

5.3 Training Stability

- Lambda weighting for different loss components
- Learning rate scheduling
- Instance normalization

6. Limitations and Future Work

6.1 Current Limitations

- Limited diversity in output styles within Monet range
- Occasional artifacts in complex scenes
- Computational intensity limiting resolution

6.2 Future Improvements

- Adaptive instance normalization for better style control
- Attention mechanisms to improve detail preservation
- Style interpolation between multiple artists
- Higher resolution support

7. Conclusion

The CycleGAN-based Monet style transfer project successfully demonstrates significant improvements over the SimpleGAN baseline. The cycle consistency constraint enables unpaired image translation while preserving structural integrity of the original photographs. The implemented optimizations address core GAN limitations, resulting in stable training and high-quality artistic renderings that capture Monet's distinctive impressionist style.

The training logs analysis confirms the stability and successful convergence of the model, with all loss components showing consistent values in the final training phase. The balance maintained between the generator and discriminator networks throughout training directly contributed to the high-quality results observed in the final model.

8. References

1. Zhu, J. Y., Park, T., Isola, P., & Efros, A. A. (2017). Unpaired image-to-image translation using cycle-consistent adversarial networks. In Proceedings of the IEEE international conference on computer vision.

2. PyTorch. (2025). PyTorch documentation — PyTorch 2.7 documentation. Retrieved from https://pytorch.org/docs/stable/index.html
3. Lindernoren, E. (2019). PyTorch-GAN: PyTorch implementations of Generative Adversarial Networks. Retrieved from https://github.com/eriklindernoren/PyTorch-GAN
4. Johnson, J., Alahi, A., & Fei-Fei, L. (2016). Perceptual losses for real-time style transfer and super-resolution. In European conference on computer vision.
5. Isola, P., Zhu, J. Y., Zhou, T., & Efros, A. A. (2017). Image-to-image translation with conditional adversarial networks. In Proceedings of the IEEE conference on computer vision and pattern recognition.